

# XGBoost Python

```
In [3]: import numpy as np
import pandas as pd
import xgboost as xgb

import matplotlib
import matplotlib as mpl
import matplotlib.pyplot as plt

from sklearn.datasets._california_housing import fetch_california_housing
from sklearn.model_selection import train_test_split
```

```
In [5]: cal_housing = fetch_california_housing(data_home="/tmp")
print(cal_housing.feature_names, cal_housing.data.shape)

['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude'] (20640, 8)
```

```
In [6]: X_train, X_test, y_train, y_test = train_test_split(cal_housing.data,
                                                         cal_housing.target,
                                                         test_size=0.4,
                                                         random_state=123)
```

```
In [7]: feature_names = ['MedInc', 'AveOccup', 'HouseAge']
feature_ids = [cal_housing.feature_names.index(f) for f in feature_names]
```

```
In [8]: dtrain = xgb.DMatrix(X_train[:, feature_ids].reshape(len(X_train), len(feature_ids)), label = y_train)
dtest = xgb.DMatrix(X_test[:, feature_ids].reshape(len(X_test), len(feature_ids)), label = y_test)
```

```
In [9]: feature_monotones = [1, -1, 1]

params = {'max_depth': 2,
          'eta': 0.1,
          'silent': 1,
          'nthread': 2,
          'seed': 0,
          'eval_metric': 'rmse',
          'monotone_constraints': '(' + ','.join([str(m) for m in feature_monotones]) + ')'
        }

bst_cv = xgb.cv(params, dtrain, 500, nfold = 5, early_stopping_rounds=10)
```

[23:38:19] WARNING: ../src/learner.cc:576:  
Parameters: { "silent" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

[23:38:19] WARNING: ../src/learner.cc:576:  
Parameters: { "silent" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

[23:38:19] WARNING: ../src/learner.cc:576:  
Parameters: { "silent" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

[23:38:19] WARNING: ../src/learner.cc:576:  
Parameters: { "silent" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

[23:38:19] WARNING: ../src/learner.cc:576:  
Parameters: { "silent" } might not be used.

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

```
In [10]: evallist = [(dtrain, 'train'), (dtest, 'eval')]

evals_result = {}
bst = xgb.train(params, dtrain, num_boost_round = bst_cv.shape[0], evals
_result = evals_result, evals = evallist, verbose_eval = False)
```

```
[23:38:44] WARNING: ../src/learner.cc:576:
Parameters: { "silent" } might not be used.
```

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

```
In [11]: print('Number of boosting rounds %d,\
              Training RMSE: %.4f, \
              Testing RMSE: %.4f' % \
              (len(evals_result['train']['rmse']),
               evals_result['train']['rmse'][-1],
               evals_result['eval']['rmse'][-1]))
```

```
Number of boosting rounds 135,          Training RMSE: 0.6954,          Test
ing RMSE: 0.6926
```

```

In [12]: def partial_dependency(bst, X, y, feature_ids = [], f_id = -1):

    """
    Calculate the dependency (or partial dependency) of a response variable
    on a predictor (or multiple predictors)
    1. Sample a grid of values of a predictor.
    2. For each value, replace every row of that predictor with this value,
    calculate the average prediction.
    """

    X_temp = X.copy()

    grid = np.linspace(np.percentile(X_temp[:, f_id], 0.1),
                       np.percentile(X_temp[:, f_id], 99.5),
                       50)
    y_pred = np.zeros(len(grid))

    if len(feature_ids) == 0 or f_id == -1:
        print('Input error!')
        return
    else:
        for i, val in enumerate(grid):

            X_temp[:, f_id] = val
            data = xgb.DMatrix( X_temp[:, feature_ids].reshape( (len(X_t
emp), len(feature_ids)) ) )

            y_pred[i] = np.average(bst.predict(data, ntree_limit = bst.b
est_ntree_limit))

    return grid, y_pred

```

```

In [13]: fig, ax = plt.subplots(1, len(feature_names))
fig.set_size_inches(len(feature_names) * 4, 5)
plt.subplots_adjust(left = 0.07, right = 0.94, bottom = 0.15, top = 0.9)

for i, f in enumerate(feature_names):

    grid, y_pred = partial_dependency(bst,
                                     X_train,
                                     y_train,
                                     feature_ids = feature_ids,
                                     f_id = cal_housing.feature_names.i
index(f)
                                     )

    if i == 0 and len(feature_names) == 1:
        axis = ax
    else:
        axis = ax[i]

    axis.plot(grid, y_pred, '-', color = 'red', linewidth = 2.5, label
='fit')
    axis.plot(X_train[:, cal_housing.feature_names.index(f)], y_train,
'o', color = 'grey', alpha = 0.01)

    axis.set_xlim(min(grid), max(grid))
    axis.set_xlabel(f, fontsize = 10)
    axis.set_ylabel('Partial Dependence', fontsize = 12)

    handles, labels = axis.get_legend_handles_labels()
    axis.legend(handles, labels, loc = 'best', fontsize = 12)
plt.savefig('w_constraint_three_feature.png')


```

```

/usr/lib/python3.6/site-packages/xgboost/core.py:108: UserWarning: ntree
e_limit is deprecated, use `iteration_range` or model slicing instead.
  UserWarning
[<matplotlib.lines.Line2D object at 0x7fb3af679eb8>]
[<matplotlib.lines.Line2D object at 0x7fb3c8c26c18>]
(0.536, 11.854603999999894)
Text(0.5, 0, 'MedInc')
Text(0, 0.5, 'Partial Dependence')
<matplotlib.legend.Legend object at 0x7fb3c8bd41d0>
/usr/lib/python3.6/site-packages/xgboost/core.py:108: UserWarning: ntree
e_limit is deprecated, use `iteration_range` or model slicing instead.
  UserWarning
[<matplotlib.lines.Line2D object at 0x7fb3c8bd4e10>]
[<matplotlib.lines.Line2D object at 0x7fb3bf9e1eb8>]
(1.2832211987226727, 6.183123919308341)
Text(0.5, 45.72222222222223, 'AveOccup')
Text(415.1928104575163, 0.5, 'Partial Dependence')
<matplotlib.legend.Legend object at 0x7fb3ce29cf60>
/usr/lib/python3.6/site-packages/xgboost/core.py:108: UserWarning: ntree
e_limit is deprecated, use `iteration_range` or model slicing instead.
  UserWarning
[<matplotlib.lines.Line2D object at 0x7fb3bf9fc320>]
[<matplotlib.lines.Line2D object at 0x7fb3c8a886d8>]
(2.0, 52.0)
Text(0.5, 45.72222222222223, 'HouseAge')
Text(783.6633986928105, 0.5, 'Partial Dependence')
<matplotlib.legend.Legend object at 0x7fb3c8a8b8d0>

```

```
In [14]: %matplotlib plt
```



## XGBoost Ray

```
In [16]: from xgboost_ray import RayDMatrix, RayParams, train
```

```
In [17]: train_x = X_train[:, feature_ids].reshape(len(X_train), len(feature_ids))
train_y = y_train
train_set = RayDMatrix(train_x, train_y)

test_x = X_test[:, feature_ids].reshape(len(X_test), len(feature_ids))
test_y = y_test
test_set = RayDMatrix(test_x, test_y)
```

```
In [18]: params
```

```
{'max_depth': 2, 'eta': 0.1, 'silent': 1, 'nthread': 2, 'seed': 0, 'eval_metric': 'rmse', 'monotone_constraints': '(1,-1,1)'}
```

```
In [19]: train_x
```

```
array([[ 4.7984      ,  2.86600496, 37.          ],
       [ 5.5714      ,  3.22791293, 34.          ],
       [ 5.1814      ,  2.86779661, 46.          ],
       ...,
       [ 5.6306      ,  3.49354376,  5.          ],
       [ 3.875       ,  1.7208589 , 44.          ],
       [ 2.5156      ,  3.56609195, 20.          ]])
```



```
In [20]: evallist = [(train_set, 'train'), (test_set, 'eval')]

# bst = xgb.train(
#     params,
#     dtrain, num_boost_round = bst_cv.shape[0], evals_result = evals_re
sult, evals = evallist, verbose_eval = False)

bst = train(
    params,
    train_set,
    evals_result=evals_result,
    evals = evallist,
    num_boost_round = bst_cv.shape[0],
#     evals=[(train_set, "train")],
    verbose_eval=False,
    ray_params=RayParams(
        num_actors=1, # Number of remote actors
        cpus_per_actor=2))
```

Usage stats collection is enabled by default for nightly wheels. To disable this, run the following command: `ray disable-usage-stats` before starting Ray. See <https://docs.ray.io/en/master/cluster/usage-stats.html> for more details.

2023-02-08 23:48:42,120 INFO worker.py:1524 -- Started a local Ray instance. View the dashboard at <http://127.0.0.1:8265>

/usr/lib/python3.6/site-packages/xgboost\_ray/main.py:439: UserWarning: `num\_actors` in `ray\_params` is smaller than 2 (1). XGBoost will NOT be distributed!

f"`num\_actors` in `ray\_params` is smaller than 2 "

2023-02-08 23:48:48,947 INFO main.py:1009 -- [RayXGBoost] Created 1 new actors (1 total actors). Waiting until actors are ready for training.

2023-02-08 23:48:53,779 INFO main.py:1054 -- [RayXGBoost] Starting XGBoost training.

(\_RemoteRayXGBoostActor pid=2316) [23:48:53] task [xgboost.ray]:1400180 40759520 got new rank 0

(\_RemoteRayXGBoostActor pid=2316) [23:48:53] WARNING: ../src/learner.cc:576:

(\_RemoteRayXGBoostActor pid=2316) Parameters: { "silent" } might not be used.

(\_RemoteRayXGBoostActor pid=2316)

(\_RemoteRayXGBoostActor pid=2316) This could be a false alarm, with some parameters getting used by language bindings but

(\_RemoteRayXGBoostActor pid=2316) then being mistakenly passed down to XGBoost core, or some parameter actually being used

(\_RemoteRayXGBoostActor pid=2316) but getting flagged wrongly here. Please open an issue if you find any such cases.

(\_RemoteRayXGBoostActor pid=2316)

(\_RemoteRayXGBoostActor pid=2316)

2023-02-08 23:48:58,553 INFO main.py:1553 -- [RayXGBoost] Finished XGBoost training on training data with total N=12,384 in 9.93 seconds (4.76 pure XGBoost training time).

```

In [21]: fig, ax = plt.subplots(1, len(feature_names))
fig.set_size_inches(len(feature_names) * 4, 5)
plt.subplots_adjust(left = 0.07, right = 0.94, bottom = 0.15, top = 0.9)

for i, f in enumerate(feature_names):

    grid, y_pred = partial_dependency(bst,
                                     X_train,
                                     y_train,
                                     feature_ids = feature_ids,
                                     f_id = cal_housing.feature_names.i
index(f)
                                     )

    if i == 0 and len(feature_names) == 1:
        axis = ax
    else:
        axis = ax[i]

    axis.plot(grid, y_pred, '-', color = 'red', linewidth = 2.5, label
='fit')
    axis.plot(X_train[:, cal_housing.feature_names.index(f)], y_train,
'o', color = 'grey', alpha = 0.01)

    axis.set_xlim(min(grid), max(grid))
    axis.set_xlabel(f, fontsize = 10)
    axis.set_ylabel('Partial Dependence', fontsize = 12)

    handles, labels = axis.get_legend_handles_labels()
    axis.legend(handles, labels, loc = 'best', fontsize = 12)
plt.savefig('w_constraint_three_feature.png')


```

```

/usr/lib/python3.6/site-packages/xgboost/core.py:108: UserWarning: ntree
e_limit is deprecated, use `iteration_range` or model slicing instead.
  UserWarning
[<matplotlib.lines.Line2D object at 0x7fb338046518>]
[<matplotlib.lines.Line2D object at 0x7fb32fd76b00>]
(0.536, 11.854603999999894)
Text(0.5, 0, 'MedInc')
Text(0, 0.5, 'Partial Dependence')
<matplotlib.legend.Legend object at 0x7fb32fd76c50>
/usr/lib/python3.6/site-packages/xgboost/core.py:108: UserWarning: ntree
e_limit is deprecated, use `iteration_range` or model slicing instead.
  UserWarning
[<matplotlib.lines.Line2D object at 0x7fb32fd76da0>]
[<matplotlib.lines.Line2D object at 0x7fb32fd55240>]
(1.2832211987226727, 6.183123919308341)
Text(0.5, 45.72222222222223, 'AveOccup')
Text(415.1928104575163, 0.5, 'Partial Dependence')
<matplotlib.legend.Legend object at 0x7fb32fd4e588>
/usr/lib/python3.6/site-packages/xgboost/core.py:108: UserWarning: ntree
e_limit is deprecated, use `iteration_range` or model slicing instead.
  UserWarning
[<matplotlib.lines.Line2D object at 0x7fb32fd5c390>]
[<matplotlib.lines.Line2D object at 0x7fb32fcee7f0>]
(2.0, 52.0)
Text(0.5, 45.72222222222223, 'HouseAge')
Text(783.6633986928105, 0.5, 'Partial Dependence')
<matplotlib.legend.Legend object at 0x7fb32fd04860>

```

In [22]: %matplotlib plt



In [ ]: